Lab Experiment 1

MTech in Applied AI
Deployments of ML Models

by

Sanjeev Kumar Pandey[*]
Faculty: Dr Ajay

Submitted December 29, 2024

[*] Student ID: 31050, Enrolment No:MT23AAI001, Email: pandey.sanjeev@yahoo.com

Table of Contents

1 Question : Hyperparameter Tuning and MLflow Integration

Objective: Showcase your skills in optimizing model performance through hyperparameter tuning and tracking experiments using MLflow.

1.1 Code

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
import pandas as pd
import mlflow.sklearn
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report


class GridSearchCVImpl:
# Step 0: Initialize
    def __init__(self):
    # a) Features and Targets
        self.X = None
        self.y = None

    # b) Define train and test
        self.X_train = self.X_test = self.y_train = self.y_test = None
        self.y_pred = None

    # c) Define the Random Forest Classifier
        self.mrfc = RandomForestClassifier()
        print("Custom Random Forest Classifier initialized.")

    # d) 1. Parameter Grid Definition:
        self.param_grid = {
            'n_estimators': [50, 100, 150],
            'max_depth': [None, 10, 20],
            'min_samples_split': [2, 5, 10]
        }

    # e) define grid_search
        self.grid_search = None

# Step 1: Load the dataset
    def load_dataset(self):
        try:
            my_df = pd.read_csv("Iris.csv")
            mapping = {'Iris-setosa': 0.0, 'Iris-versicolor': 1.0, 'Iris-virginica': 2.0}
            my_df['Species'] = my_df['Species'].map(mapping).astype(float)

            # X is feature so dropping the Y from the dataframe
            # ensure to take only 4 features as model is accepting only 4 features
            self.X = my_df[['SepalLengthCm', 'SepalWidthCm', 'PetalLengthCm', 'PetalWidthCm']]
            self.y = my_df['Species']
            print("Dataset loaded successfully.")
        except Exception as e:
            print(f"Error loading dataset: {e}")
            exit()

# Step 2: Split the dataset
    def split_dataset(self):
        try:
            X = self.X.values
            y = self.y.values
            self.X_train, self.X_test, self.y_train, self.y_test = (
                train_test_split(X, y, test_size=0.2, random_state=42)
```

```
            )
            print("Dataset split into train and test sets.")
        except Exception as e:
            print(f"Error splitting dataset: {e}")
            exit()

# Step 3: Set up and run GridSearchCV
    def execute_gridsearchcv(self):
        try:
            print("Starting GridSearchCV...")
            self.grid_search = GridSearchCV(
                estimator=self.mrfc,
                param_grid=self.param_grid,
                scoring='accuracy',
                cv=5,  # 5-fold cross-validation    2. Grid Search Implementation:
                n_jobs=-1,  # Use all CPU cores
                verbose=1   # Display progress
            )
            self.grid_search.fit(self.X_train, self.y_train)
            print("GridSearchCV completed.")
        except Exception as e:
            print(f"Error during GridSearchCV: {e}")
            exit()

# Step 4: Get the best hyperparameters and evaluate
    def eval_and_report(self):
        try:
            print(f"Best Hyperparameters:{self.grid_search.best_params_}")
            print(f"Best Cross-Validation Score:{self.grid_search.best_score_}")

            # Evaluate on the test set
            best_model = self.grid_search.best_estimator_
            self.y_pred = best_model.predict(self.X_test)

            print(f"Actual : Predicted")
            for i in range(len(self.y_test)):
                print(f"{self.y_test[i]}     : {self.y_pred[i]}")
            accuracy = accuracy_score(self.y_test, self.y_pred)
            print("Test Set Accuracy:", accuracy)

            # Print classification report
            print("Classification Report:")
            print(classification_report(self.y_test, self.y_pred))
        except Exception as e:
            print(f"Error during evaluation: {e}")
            exit()


# Set up experiment tracking
mlflow.set_experiment('RandomForest_Hyperparameter_Tuning')

# Start an MLflow run
with mlflow.start_run():
    gscvi = GridSearchCVImpl()
    gscvi.load_dataset()
    gscvi.split_dataset()

    # Log hyperparameter grid
    mlflow.log_params({
        "n_estimators_range": gscvi.param_grid['n_estimators'],
        "max_depth_range": gscvi.param_grid['max_depth'],
        "min_samples_split_range": gscvi.param_grid['min_samples_split']
    })

    # Execute GridSearchCV
```

```
    gscvi.execute_gridsearchcv()

    # Log the best hyperparameters and the best score
    mlflow.log_params(gscvi.grid_search.best_params_)
    mlflow.log_metric("best_cv_accuracy", gscvi.grid_search.best_score_)

    # Log cross-validation scores
    for i, score in enumerate(gscvi.grid_search.cv_results_['mean_test_score']):
        mlflow.log_metric(f'cv_score_{i + 1}', score)
        print(f"cv_score_{i + 1}, {score}")

    # Log the best model
    input_example = gscvi.X_train[0].reshape(1, -1)  # Example input for reproducibility
    mlflow.sklearn.log_model(gscvi.grid_search.best_estimator_, "best_random_forest_model",
                             input_example=input_example)

    # Log test accuracy
    gscvi.eval_and_report()
    test_accuracy = accuracy_score(gscvi.y_test, gscvi.y_pred)
    mlflow.log_metric("test_accuracy", test_accuracy)
```

Listing 1: Code

1.2 Output

```
Custom Random Forest Classifier initialized.
Dataset loaded successfully.
Dataset split into train and test sets.
Starting GridSearchCV...
Fitting 5 folds for each of 27 candidates, totalling 135 fits
GridSearchCV completed.
cv_score_1, 0.95
cv_score_2, 0.9416666666666667
cv_score_3, 0.95
cv_score_4, 0.95
cv_score_5, 0.95
cv_score_6, 0.95
cv_score_7, 0.95
cv_score_8, 0.95
cv_score_9, 0.95
cv_score_10, 0.95
cv_score_11, 0.9333333333333333
cv_score_12, 0.95
cv_score_13, 0.95
cv_score_14, 0.95
cv_score_15, 0.95
cv_score_16, 0.95
cv_score_17, 0.95
cv_score_18, 0.95
cv_score_19, 0.9416666666666667
cv_score_20, 0.95
cv_score_21, 0.95
cv_score_22, 0.95
cv_score_23, 0.9416666666666667
cv_score_24, 0.95
cv_score_25, 0.95
cv_score_26, 0.95
cv_score_27, 0.9583333333333334
```

Figure 1: op-1

1.3 Output

Figure 2: op-2

1.4 Output

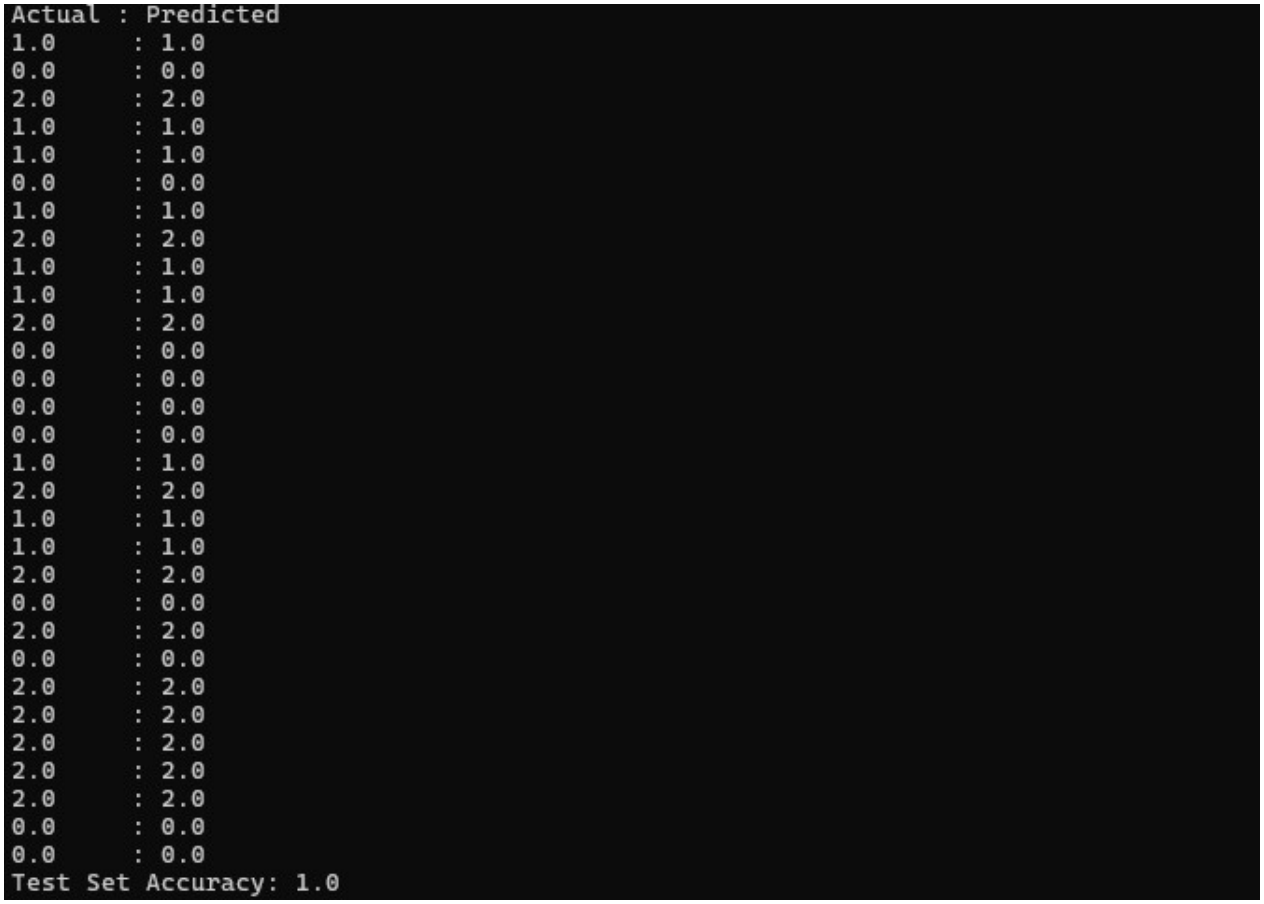

Figure 3: op-3

1.5 Output

```
Classification Report:
              precision    recall  f1-score   support

         0.0       1.00      1.00      1.00        10
         1.0       1.00      1.00      1.00         9
         2.0       1.00      1.00      1.00        11

    accuracy                           1.00        30
   macro avg       1.00      1.00      1.00        30
weighted avg       1.00      1.00      1.00        30
```

Figure 4: op-4

1.6 Output

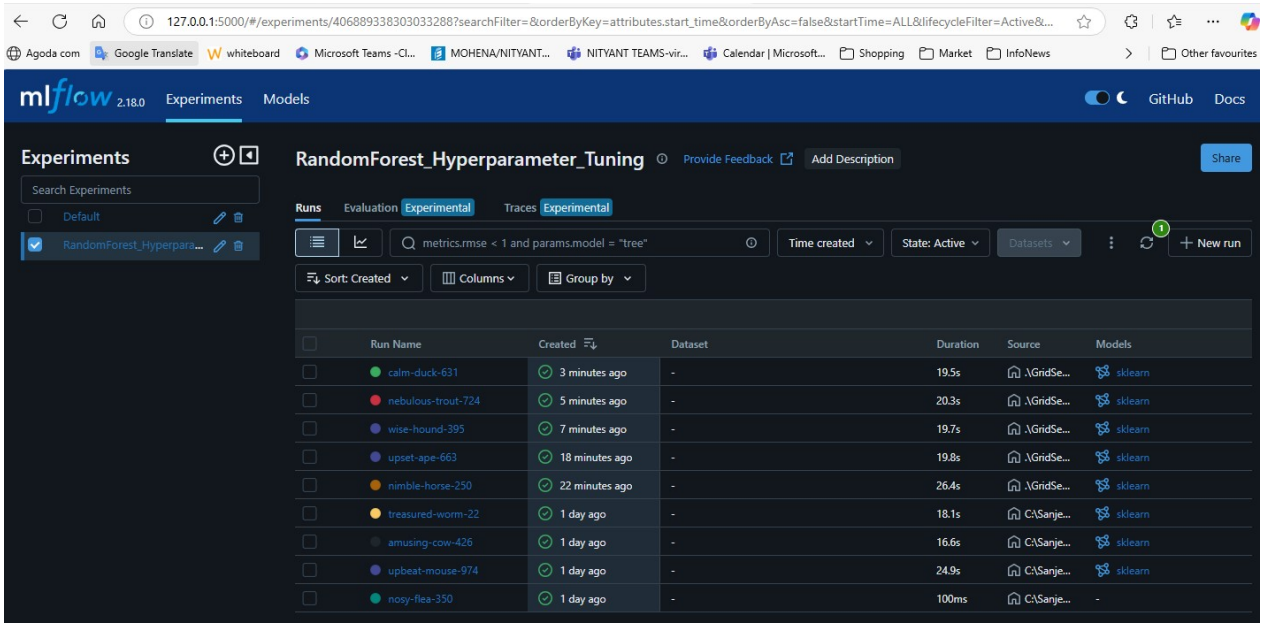

Figure 5: op-5

1.7 Output

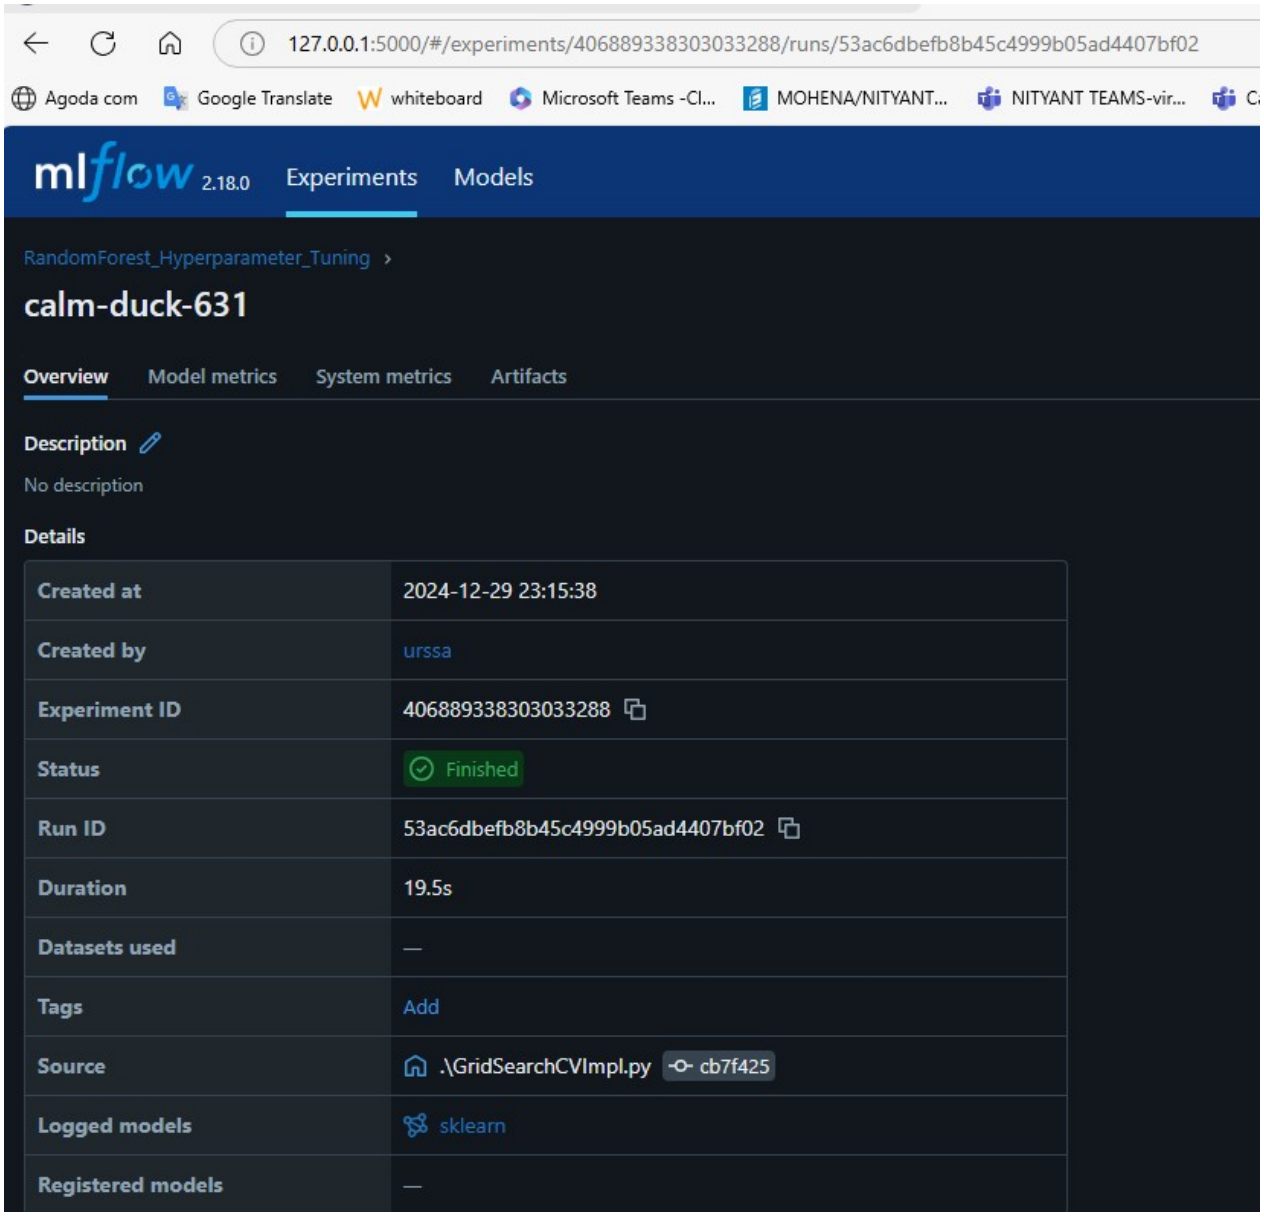

Figure 6: op-6

1.8 Output

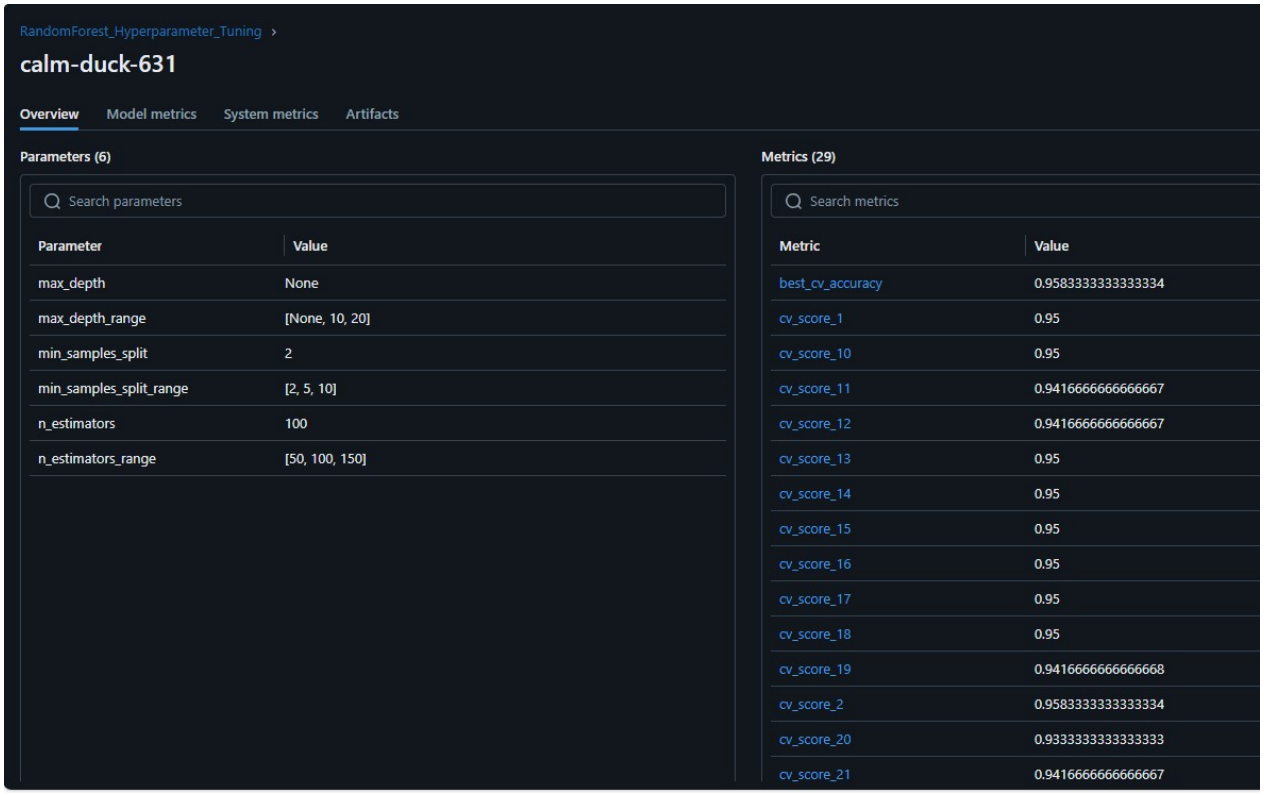

Figure 7: op-7

1.9 Output

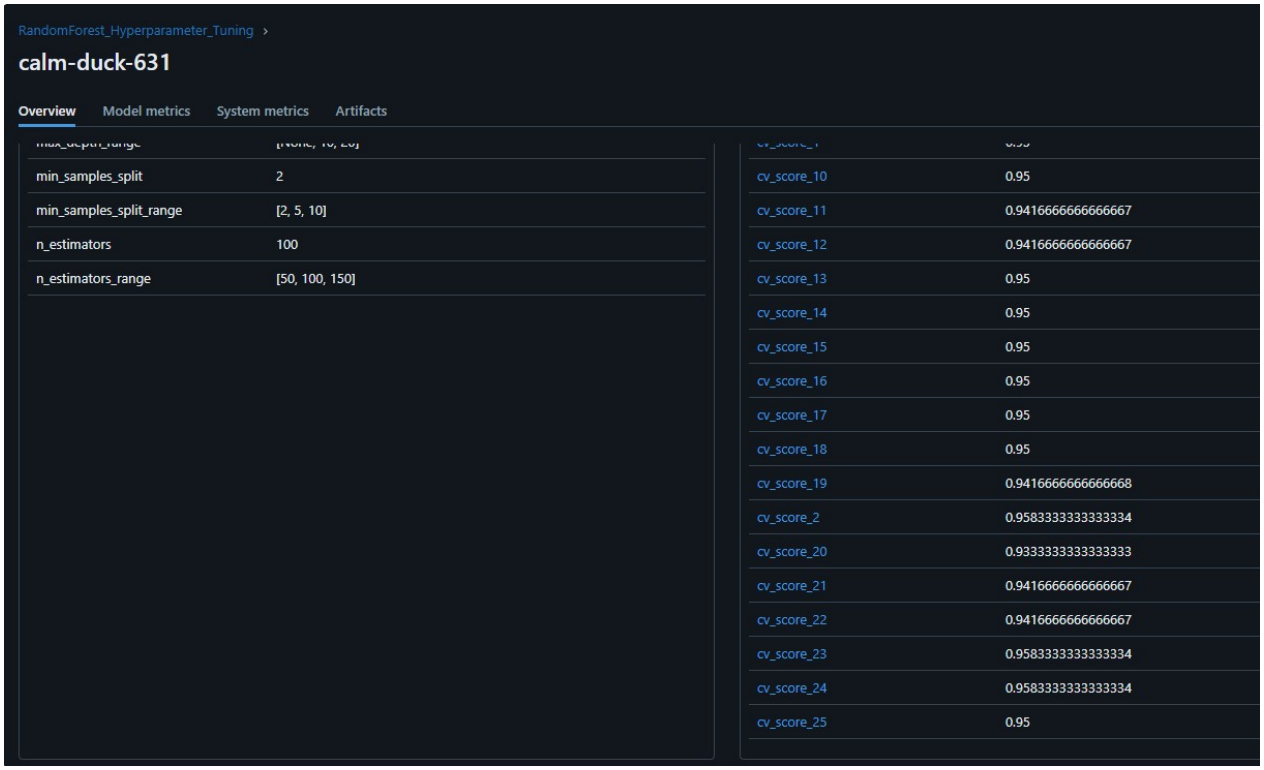

Figure 8: op-8

1.10 Output

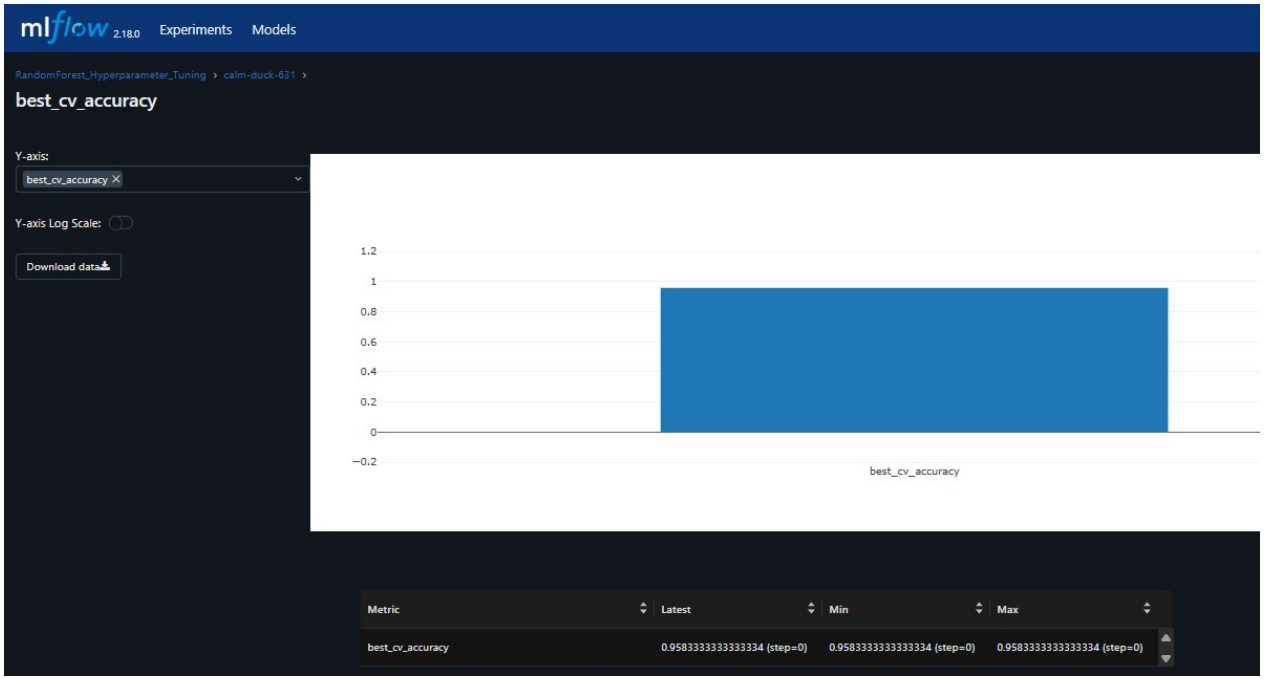

Figure 9: op-9

1.11 Output

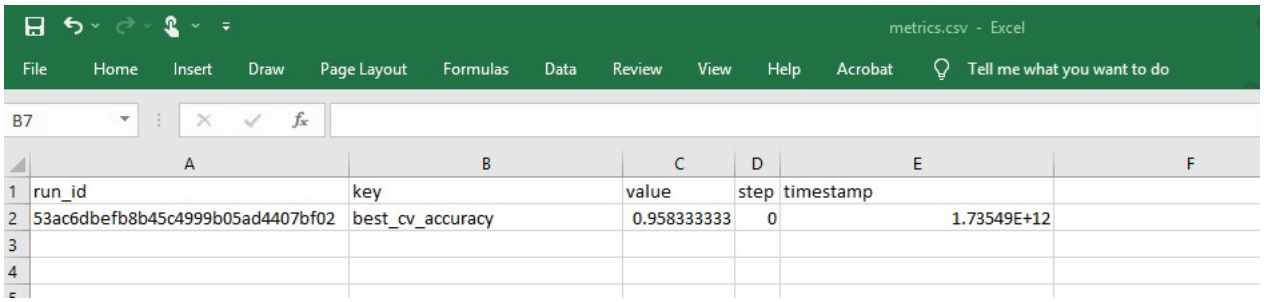

Figure 10: op-10

1.12 Output

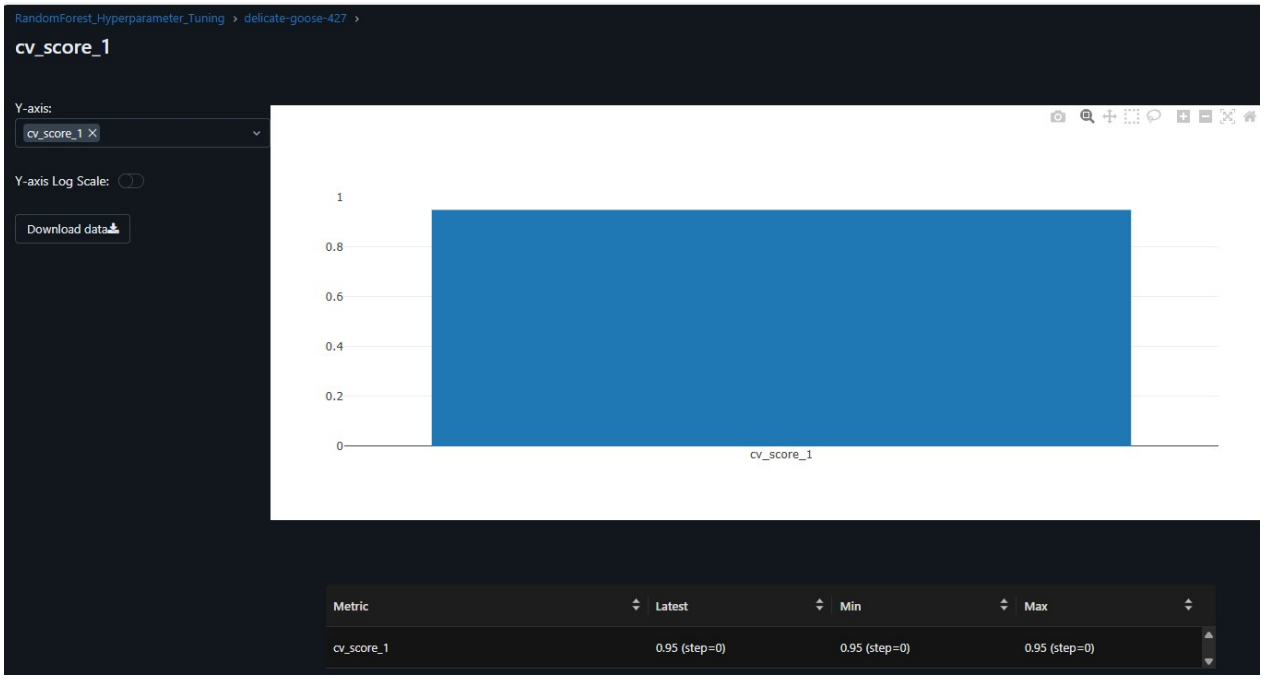

Figure 11: op-11

1.13 Output

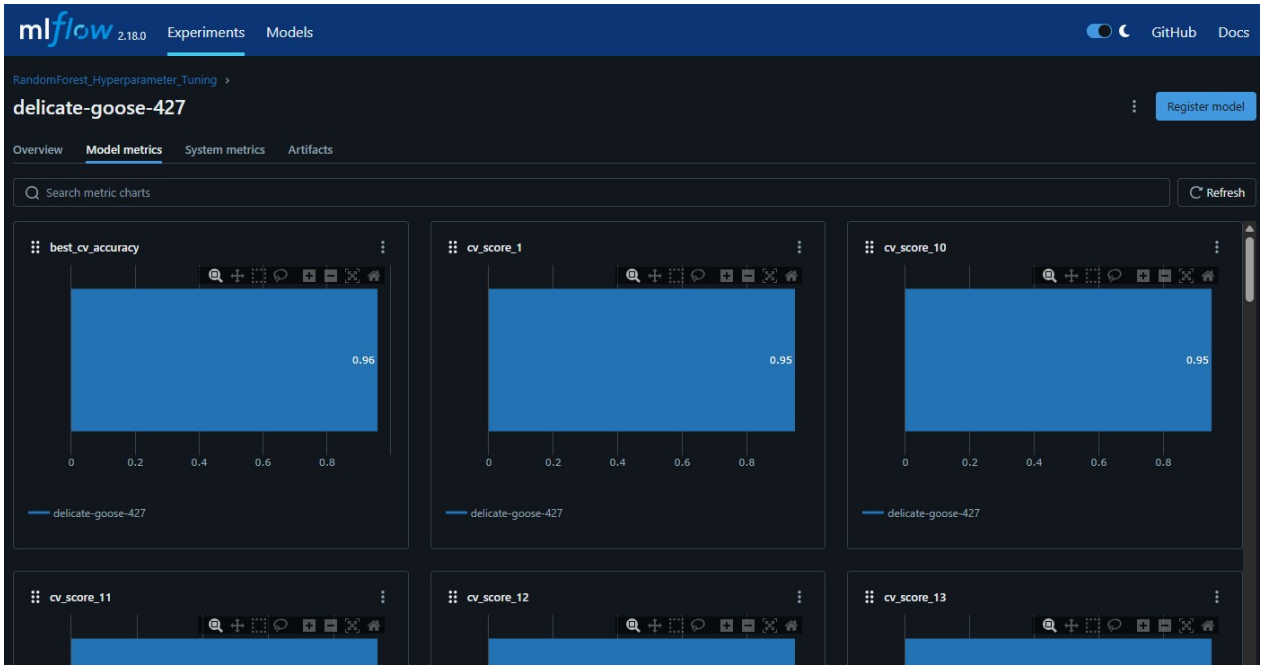

Figure 12: op-12

1.14 Output

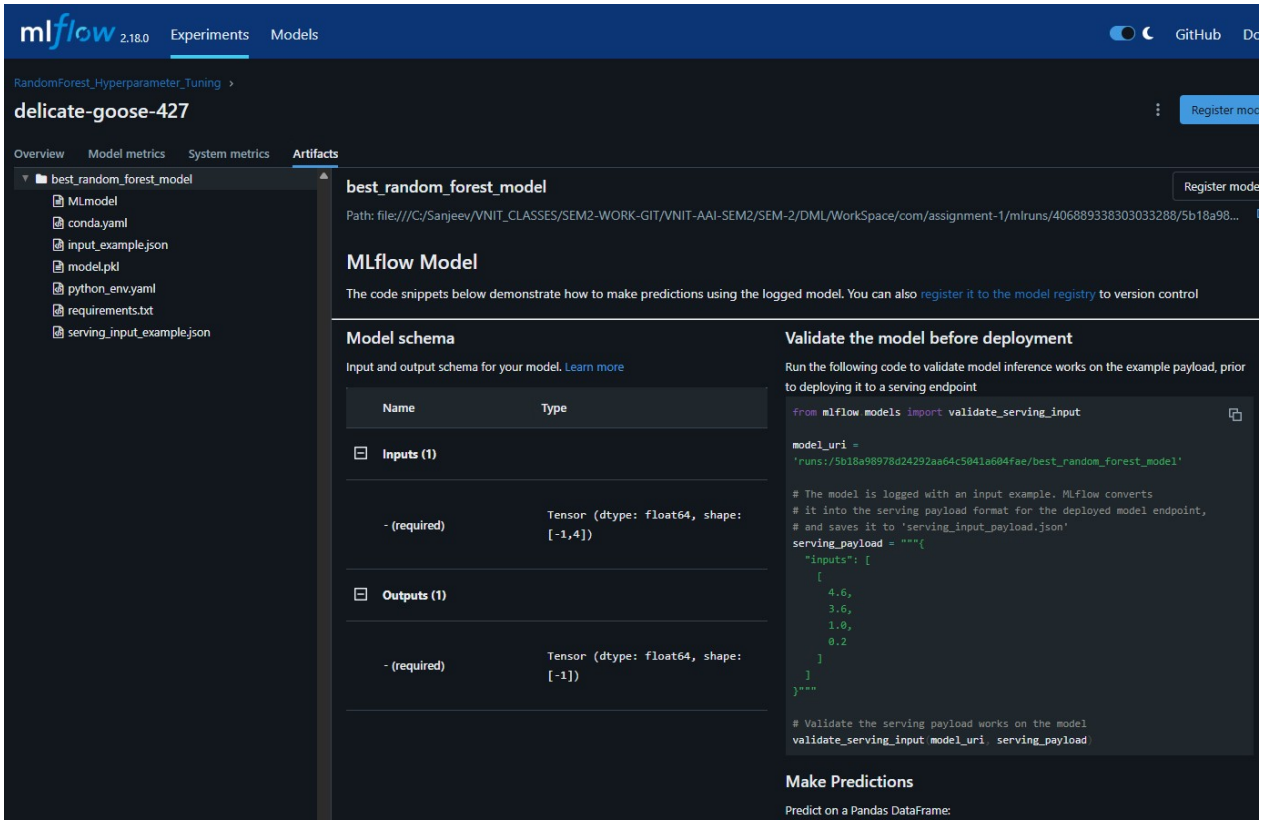

Figure 13: op-13

1.15 Output

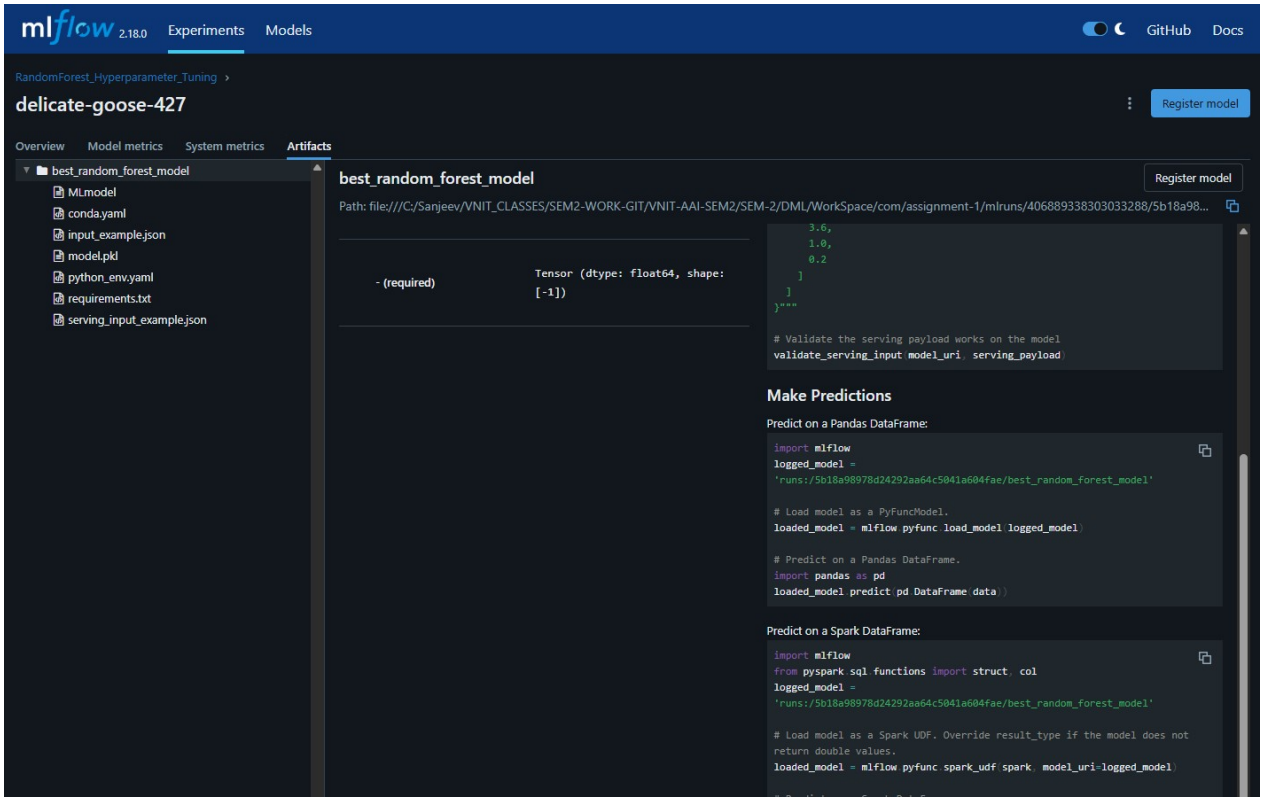

Figure 14: op-14

1.16 Output

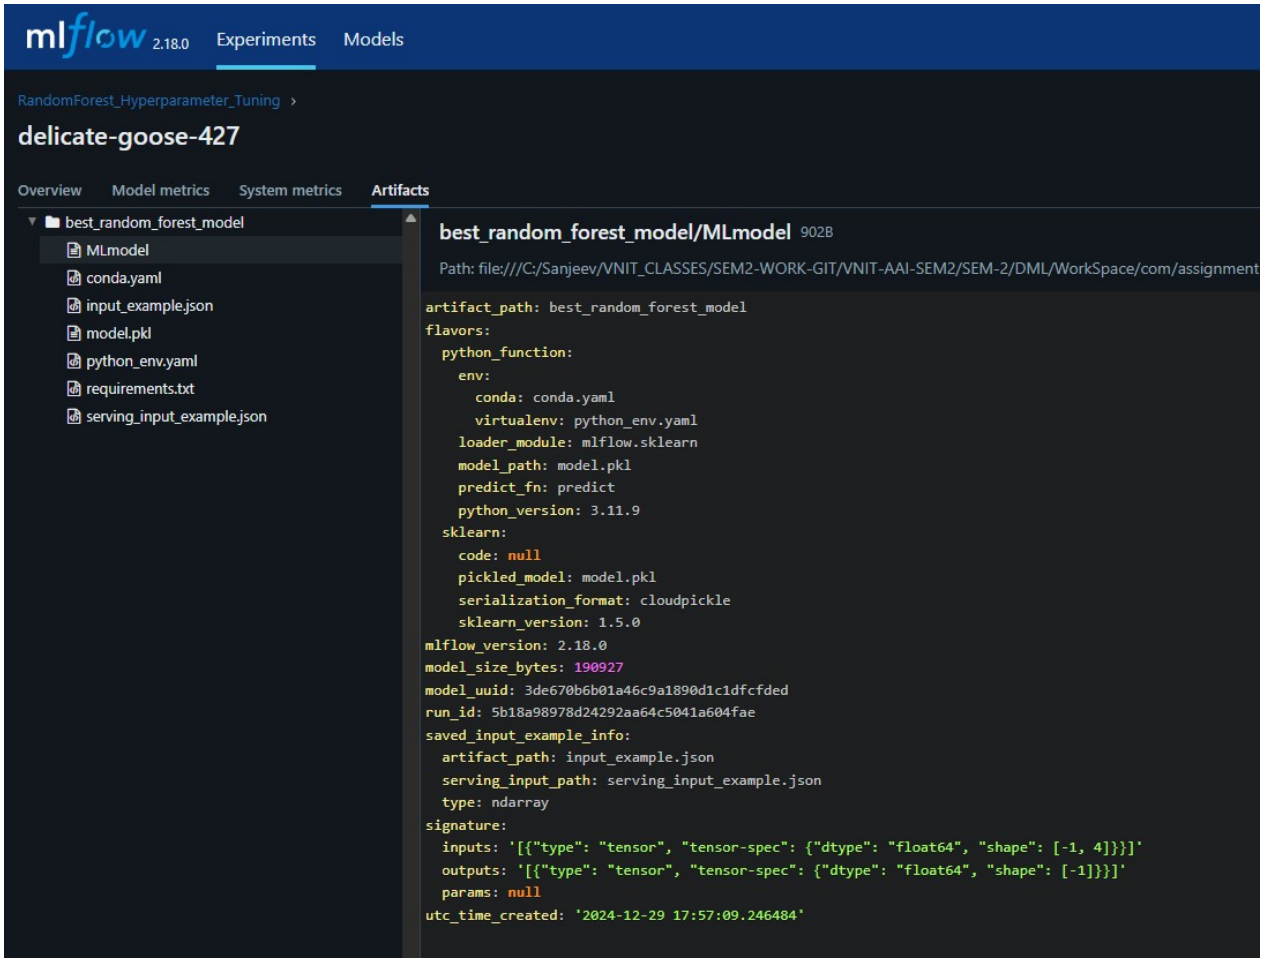

Figure 15: op-15

1.17 Output

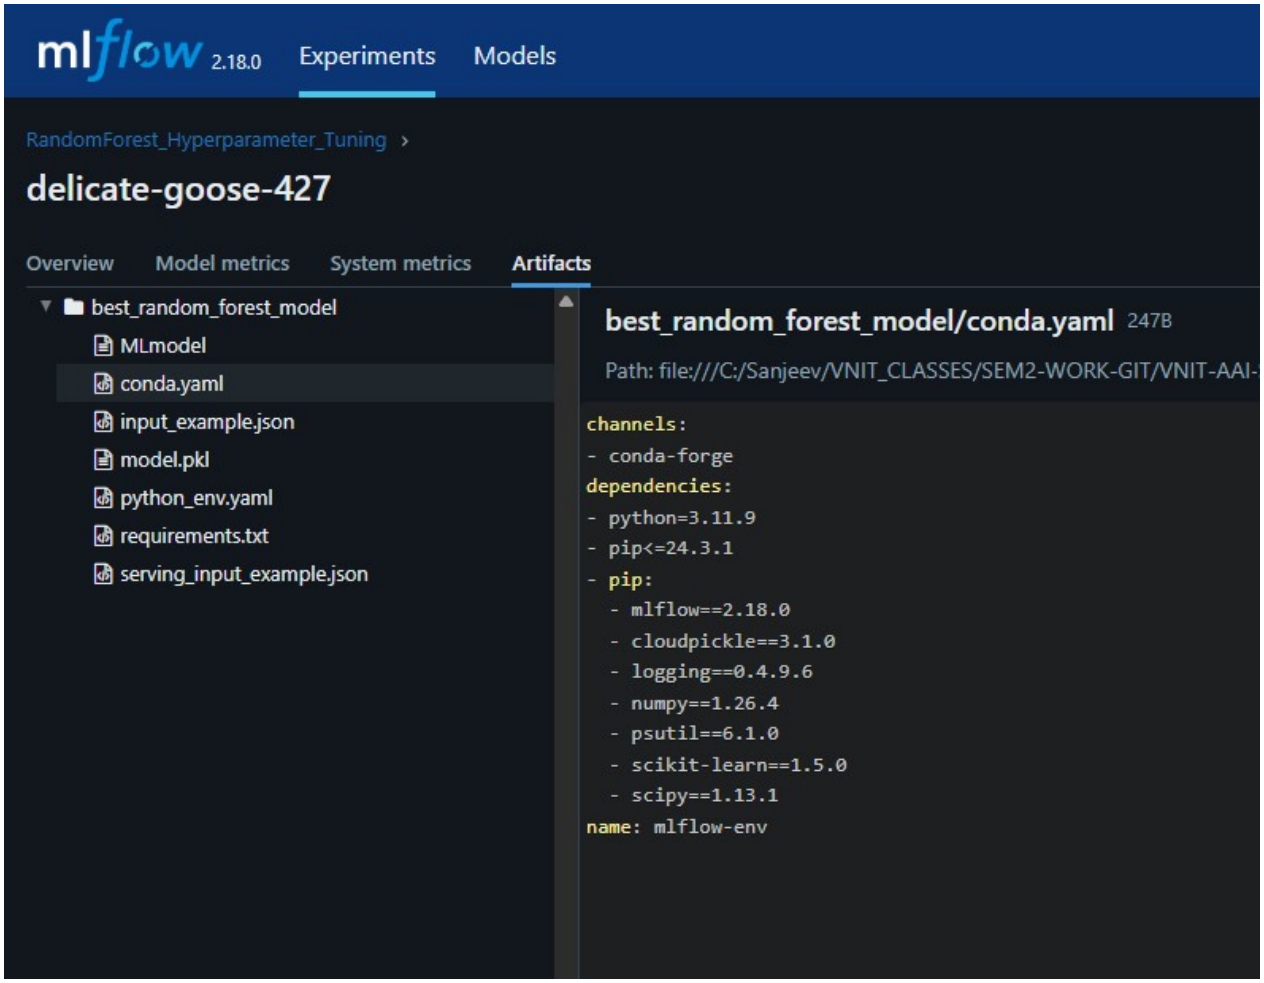

Figure 16: op-16

1.18 Output

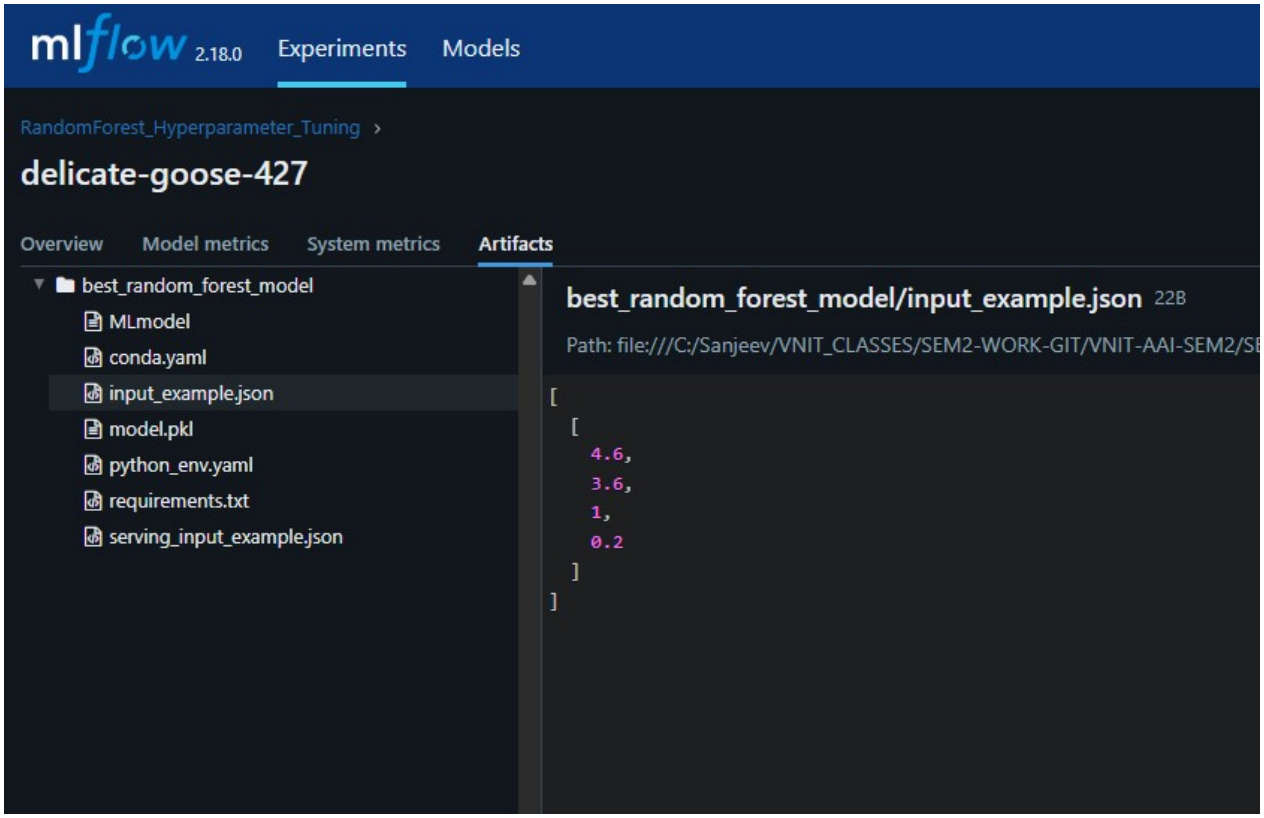

Figure 17: op-17

1.19 Output

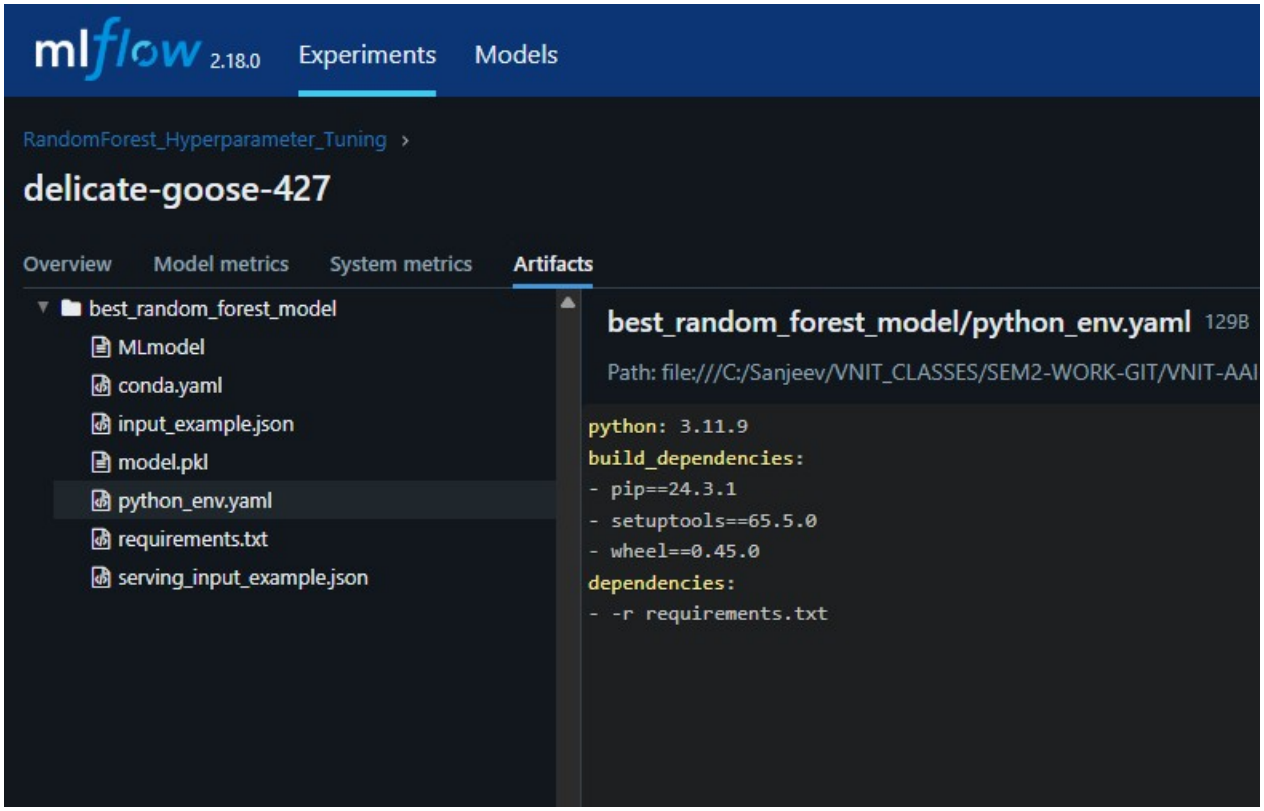

Figure 18: op-18

1.20 Output

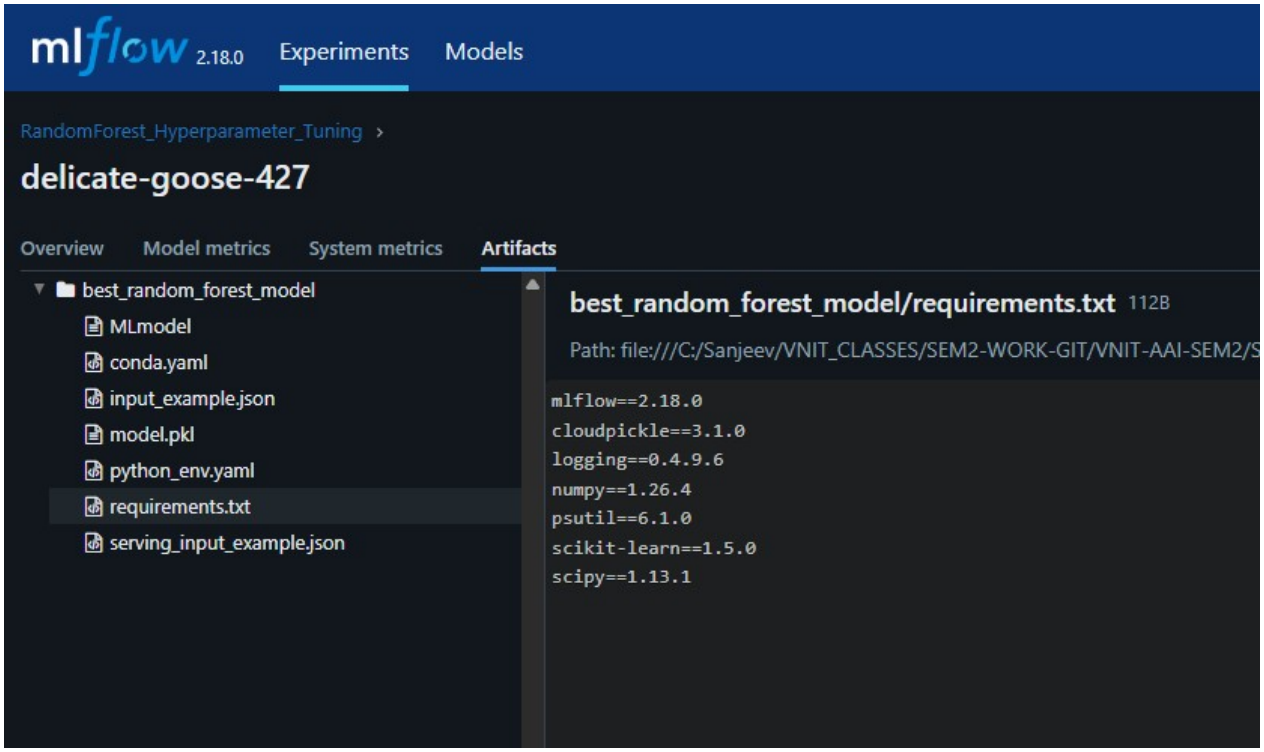

Figure 19: op-19

1.21 Output

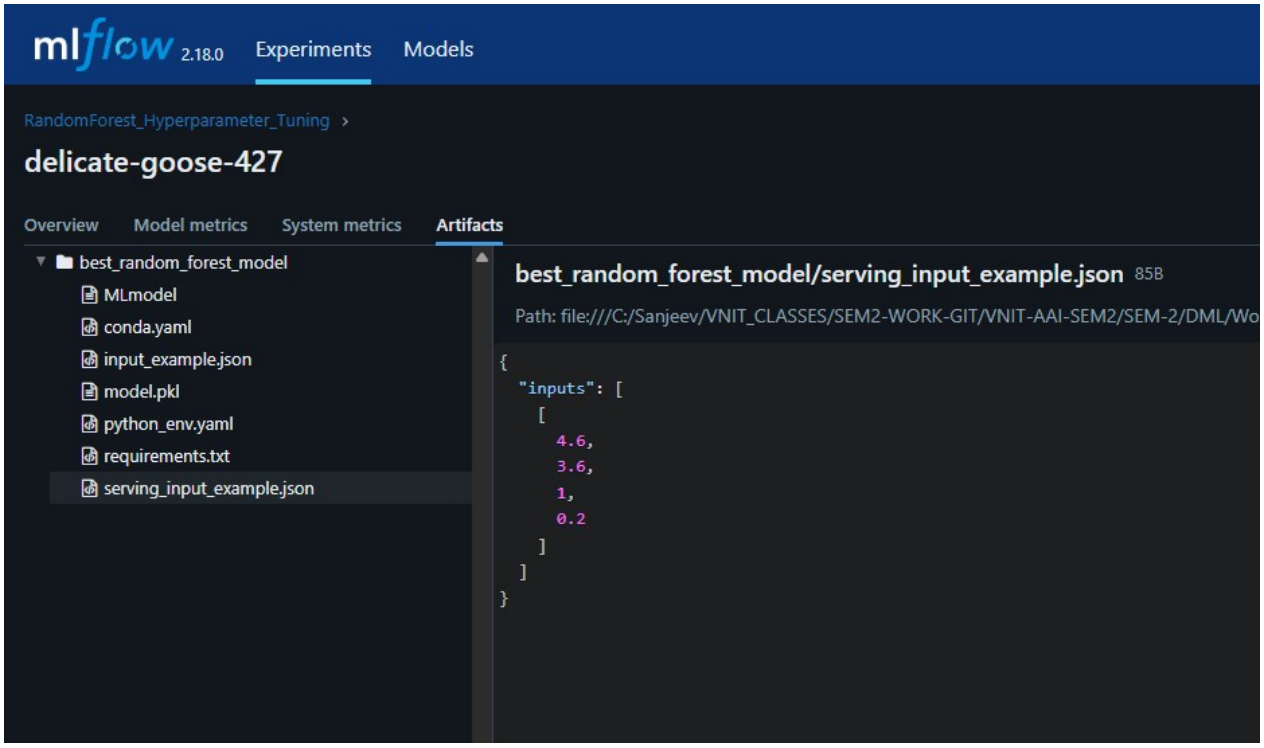

Figure 20: op-20